Table 5. Key Quantitative Findings

Measure	Observed / model value	Baseline / contrast	Interpretation
PCS matching rate	97.02%	Random baseline 5.07%	Structured recombination reproduces observed forms far better than random generation.
Token mass coverage	85.25%	Random baseline 18.50%	Frequent corpus mass is captured by the PCS model.
Zipf slope difference	0.0795	Random baseline 0.7138	Generated distribution remains close to real frequency structure.
H(suffix|core)	0.8370	H(core|prefix) = 5.3205	Suffix selection is much more constrained by core than core is by prefix.
Compression ratio	Real 0.3089; PCS 0.3163	Shuffled 0.3240	Real and PCS-generated corpora are more compressible than shuffled text.
Bigram entropy	Real 11.7001	Line-internal shuffle 12.7229	Real local ordering is more constrained than line-preserving randomization.
Next-token prediction	Top-1 0.6917; top-3 0.8168	Baseline top-1 0.5505; top-3 0.6912	Observed local sequences improve next-token predictability.
Line-position dependence	Chi-square = 46151.0531	$p < .001$	Line position is statistically associated with token distribution.
Family clustering	Same-core mean distance 235.93	Shuffled 442.77	Related token families cluster locally.

Table 5. Key quantitative findings summary. Values are drawn from the component studies and are used as cross-level constraints for future models.